

Hypothesis Evidence Landscape and Temporal Dynamics

The LLM-based extraction pipeline produced a total of $\{\{\text{TOTAL_ASSERTIONS}\}\}$ assertions across the eight tracked hypotheses, drawn from the full corpus of $N = \text{CORPUS_SIZE}$ papers. The distribution of assertion types and the resulting citation-weighted scores reveal a differentiated evidence landscape:

Hypothesis	Score	Supports	Neutral	Contradicts	Total	Character
H4: Predictive Coding	+0.59	837	417	0	1,254	Strong consensus
H5: Scalability	+0.62	142	110	0	252	Strong consensus
H6: Clinical Utility	+0.41	16	29	0	45	Moderate, growing
H8:	+0.39	54	96	0	150	Moderate,